

# **10 Reasons why you should prefer PostgreSQL to MySQL**

Anand Chitipothu

# **Who am I?**

## **Anand Chitipothu**


Independent Software Consultant & Trainer

Ex. Internet Archivist

# MySQL or PostgreSQL?



The world's most **popular**  
open source database



The world's most **advanced**  
open source database

# Quick Comparison

	<b>MySQL MyISAM</b>	<b>MySQL InnoDB</b>	<b>PostgreSQL</b>
Transactions	No	Yes	Yes
Foreign Key Constraints	No	Yes	Yes
Locking	Table	Row/MVCC	Row/MVCC

# Who Uses MySQL?

- Facebook
- FriendFeed
- Quora
- Flickr
- Second Life

# Who Uses PostgreSQL

- Skype
- Heroku
- Instagram
- Disqus
- Yahoo!
- NASA

**MySQL ate my cake**

# MySQL ate my cake

```
mysql> CREATE TABLE cake (name VARCHAR(3));
```

```
Query OK, 0 rows affected (0.02 sec)
```

```
mysql> INSERT INTO cake (name) VALUES  
('pancake');
```

```
Query OK, 1 row affected, 1 warning (0.03 sec)
```



# MySQL ate my cake

```
mysql> SELECT * FROM cake;
```

```
+-----+
```

```
| name |
```

```
+-----+
```

```
| pan  |
```

```
+-----+
```

```
1 row in set (0.03 sec)
```

**OMG! Where is my “cake”?**

# PostgreSQL?

```
testdb=# CREATE TABLE cake (name VARCHAR(3));  
CREATE TABLE
```

```
testdb=# INSERT INTO cake (name)  
VALUES ('pancake');
```

```
ERROR:  value too long for type character  
varying(3)
```

# Ever Seen This?

```
CREATE TABLE users (  
    id integer auto_increment,  
    username varchar(40),  
    password varchar(8)  
);
```

# Data Conversion Errors - MySQL

```
mysql> CREATE TABLE test (x INTEGER);
```

```
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> INSERT INTO test (x)  
VALUES ('bad-number');
```

```
Query OK, 1 row affected, 1 warning (0.01 sec)
```

# Data Conversion Errors - MySQL

```
mysql> SELECT * FROM foo;
```

```
+-----+
```

```
| x      |
```

```
+-----+
```

```
|      0 |
```

```
+-----+
```

```
1 row in set (0.00 sec)
```

# Data Conversion Errors - PostgreSQL

```
testdb=# CREATE TABLE test (x integer);
```

```
CREATE TABLE
```

```
testdb=# INSERT INTO test (x) VALUES ('bad-  
number');
```

```
ERROR:  invalid input syntax for integer: "bad-number"
```

```
LINE 1: INSERT INTO foo (x) VALUES ('bad-number');
```

^

# Parallels - PHP vs. Python

```
$ php -r '$x = "bad-number";  
          $y = (int)$x;  
          echo $y."\n";'
```

0

# Parallels - PHP vs. Python

```
$ python -c 'print int("bad-number")'  
Traceback (most recent call last):  
  File "<string>", line 1, in <module>  
ValueError: invalid literal for int()  
with base 10: 'bad-number'
```



# **File Layout**

# File Layout - MySQL MyISAM

`/var/lib/mysql/dbname`

- `dbname.MYD` - data of all tables
- `dbname.MYI` - indexes

# File Layout - MySQL InnoDB

**/var/lib/mysql/**

- **ibdata1** - data of all databases, including tables and indexes

It is possible to tell mysql, by changing a config flag, to make it use one file for table.

# File Layout - PostgreSQL

**/var/lib/postgresql/9.3/main/base**

- **131384/** - directory per database
  - **2654** - one (or more) files for each table/index
  - **2703**
  - **2683**
  - **2683.1**
  - **2683.2**
  - **...**

# **Database Maintenance**

# CREATE INDEX - MySQL MyISAM

While CREATE INDEX is in progress:

- Entire table is locked for writes
- A new index file (dbname.MYI) need to created

# CREATE INDEX - InnoDB

- Entire table rebuilt to create an index.
- Seems to have improved in recent versions

# CREATE INDEX - PostgreSQL

- **CREATE INDEX**
  - locks the table for writes
- **CREATE INDEX CONCURRENTLY**
  - Doesn't hold the lock for entire period
  - Slower than plain CREATE INDEX
- Each index is a new file



# DROP INDEX - MySQL

Takes long time as it needs to rewrite:

- the index file (**dbname.MYI**) for MyISAM
- the **ibdata1** file for InnoDB

# DROP INDEX - PostgreSQL

- Almost instantaneous
- Just need to delete the files corresponding to that index

# **ADDING NEW COLUMN - MySQL**

Entire table data needs to be rewritten.

# ADDING NEW COLUMN - PostgreSQL

Almost instantaneous if the new column has a default value.

# Connection Model

# Connection Model - MySQL

A thread for each connection

## **PROS**

- Very easy to create a new conn

## **CONS**

- Difficult to scale on multi-core systems
- difficult monitor threads

# Connection Model - PostgreSQL

A process for each connection

## **PROS**

- better concurrency
- complete isolation
- plays nicely with UNIX tools (ps, top, kill)

## **CONS**

- lot of overhead for creating new conn

\$ top

PID	%CPU	%MEM	COMMAND	
17618	82.7	6.8	postgres: anand voterdb [local] COPY	+
11894	0.0	7.1	postgres: [redacted] [local] idle	+
16091	0.0	2.2	/usr/lib/postgresql/9.3/bin/postgres -D /va	+
16093	0.0	17.3	postgres: checkpoint process	+
16094	0.0	0.4	postgres: writer process	+
16095	0.0	0.9	postgres: wal writer process	+
16096	0.0	0.1	postgres: autovacuum launcher process	+
16097	0.0	0.1	postgres: stats collector process	+
30283	0.0	6.7	postgres: [redacted] [local] idle	+



\$ top

PID	%CPU	%MEM	COMMAND	
17618	82.7	6.8	postgres: anand voterdb [local] COPY	+
11894	0.0	7.1	postgres: [REDACTED] [local] idl+	
16091	0.0	2.2	/usr/lib/postgresql/9.3/bin/postgres -D /va+	
16093	0.0	17.3	postgres: checkpointer process	+
16094	0.0	0.4	postgres: writer process	+
16095	0.0	0.9	postgres: wal writer process	+
16096	0.0	0.1	postgres: autovacuum launcher process	+
16097	0.0	0.1	postgres: stats collector process	+
30283	0.0	6.7	postgres: [REDACTED] [local] idl+	

**kill 17618**

**kill -STOP 17618**

**kill -CONT 17618**

# Query Planning

# The Query

**EXPLAIN**

**SELECT** name, total

**FROM** names

**WHERE** year=1995

**ORDER BY** total **DESC**

**LIMIT** 10;

# MySQL - Query Plan

id	select_type	table	type	possible_keys
1	SIMPLE	names	ALL	NULL

key	key_len	ref	rows	Extra
NULL	5	const	12420176	Using where;
				Using filesort

1 row in set (0.00 sec)

# MySQL - Add Index

```
CREATE INDEX names_total_idx  
ON names(total)
```

# MySQL - Query Plan With Index

id	select_type	table	type	possible_keys
1	SIMPLE	names	ref	NULL
key	key_len	ref	rows	Extra
names_total_idx	5	const	10	Using where;

1 row in set (0.00 sec)

# PostgreSQL - QUERY PLAN

**Limit** (cost=246987.39..246987.42 rows=10 width=13)

-> **Sort** (cost=246987.39..247467.64 rows=192099 width=13)

Sort Key: total

-> **Seq Scan on names**

(cost=0.00..242836.20 rows=192099 width=13)

Filter: (year = 1995)

(5 rows)

# PostgreSQL - Add Index

```
CREATE INDEX names_total_idx  
ON names(total)
```



# PostgreSQL - Query Plan With Index

**Limit** (cost=0.43..1891.80 rows=10 width=13)

-> **Index Scan Backward** using  
names\_total\_idx on names

(cost=0.43..36332875.67 rows=192099 width=13)

**Filter:** (year = 1995)

(3 rows)

# PostgreSQL - Query Plan

names=# **EXPLAIN ANALYZE** SELECT ...

Limit (cost=0.43..1891.80 rows=10 width=13)

(actual time=0.037..0.462 rows=10 loops=1)

-> Index Scan Backward using names\_total\_idx  
on names (cost=0.43..36332875.67 rows=192099 width=13)  
(actual time=0.036..0.458 rows=10 loops=1)

Filter: (year = 1995)

Rows Removed by Filter: 467

Total runtime: 0.517 ms

(5 rows)

# PostgreSQL - Complex Query

**EXPLAIN**

```
SELECT name, sum(total) as count  
FROM names  
WHERE year > 1980  
GROUP BY name  
ORDER BY count
```

# PostgreSQL - Query Plan

Sort (cost=254889.31..255063.44 rows=69653 width=13)

Sort Key: (sum(total))

-> HashAggregate(cost=248589.93..249286.46 rows=69653 width=13)

-> Bitmap Heap Scan on names

(cost=83212.02..226363.10 rows=4445366 width=13)

Recheck Cond: (year > 1980)

-> Bitmap Index Scan on names\_year\_idx

(cost=0.00..82100.68 rows=4445366 width=0)

Index Cond: (year > 1980)

(7 rows)

# Replication

# MySQL Replication

## Replication Format


- Statement based
- Row based
- Mixed

## Mode

- binlog
- GTID

# PostgreSQL Replication

- Synchronous / Asynchronous
- Streaming / Log shipping



# **Data Recovery**



# PG - Point In Time Recovery

- Postgres maintains Write Ahead Log of all changes made to the database.
- The WAL files can be replayed up to any given timestamp.

**Time machine of your database!**

# **Other Interesting Features of PostgreSQL**

# Partial Indexes

```
SELECT * FROM comments  
WHERE email LIKE '%@spam.com';
```

```
CREATE INDEX comments_spam_idx  
ON comments  
WHERE email LIKE '%@spam.com';
```

# Functional Indexes

```
SELECT tag, count(*) FROM posts  
GROUP BY lower(category);
```

```
CREATE INDEX post_category_idx  
ON post (lower(category));
```

# JSON datatype

```
CREATE TABLE book (  
    id serial primary key,  
    data JSON  
);
```

```
INSERT INTO book (data) VALUES (  
    '{"title": "Tom Sawyer",  
    "author": "Mark Twain"}')
```

# JSON datatype

```
SELECT * FROM book
WHERE data->>'author' = 'Mark Twain'
```

id	data
-----	
1	{"title": "Tom Sawyer", "author": "Mark Twain"}

(1 row)

# pg\_stat\_statements

```
SELECT total_time/calls AS t, calls, query  
FROM pg_stat_statements  
ORDER BY t DESC
```

```
8606.75|3|select name, sum(total) as count  
      | |from names group by name order by count limit ?;  
4.92|8| select name, total from names  
      | | where year=? order by total desc limit ?;
```

# Summary

PostgreSQL is better than MySQL in

- Data Consistency
- Query Planning
- Stability
- Database Maintenance
- Data Recovery

Worth trying PostgreSQL for your next project!



# Credits

PostgreSQL Logo - By Jeff MacDonald

[http://pgfoundry.org/docman/?group\\_id=1000089](http://pgfoundry.org/docman/?group_id=1000089)

# Thanks!

Anand Chitipothu  
@anandology