

Logical Replication in PostgreSQL

Petr Jelínek
PostgreSQL Contributor & Consultant



History



Postgres Replication History

- PostgreSQL 7.x (~2000)
 - Replication should not be part of core Postgres
- PostgreSQL 8.0 (2005)
 - Point-In-Time Recovery
- PostgreSQL 9.0 (2010)
 - Streaming Replication (physical)
- PostgreSQL 9.4 (2014)
 - Logical Decoding (changeset extraction)



Logical Replication History

- Trigger based solutions
 - Slony (~2004)
 - Londiste (~2007)
- Result of the no replication in core philosophy
- Use table(s) as queue
 - Amplify writes on the upstream
- Written in scripting languages
- But work for the most part



BDR Project

- Native logical replication
- Multi-master
- Firm design 2011
- Prototype 2012
- Resulted in many patches to PostgreSQL 9.3, 9.4, 9.5 and now 9.6
 - For example Logical Decoding
- Not **yet** fully integrated into PostgreSQL



Current state


- BDR
 - Modified PostgreSQL 9.4 + Extension
 - Multi-master
 - Transparent DDL
- pglogical
 - Extension for 9.4+
 - One way replication
 - Replacement for trigger-based solutions
 - Integrating into PostgreSQL core



Streaming Replication




Physical Streaming Replication





Logical Streaming Replication






Logical Decoding



Logical Decoding

Master





Logical Decoding

- Extracts information from Write-Ahead Log into logical changes (INSERT/UPDATE/DELETE)
- Per row and commit ordered
- C API for output plugin
- No DDL
- SQL Interface
- Streaming Interface



Existing output plugins


- <https://github.com/confluentinc/bottledwater-pg>
 - Decodes change into AVRO
 - Sends data to Kafka
- <https://github.com/xstevens/decoderbufs>
 - Decodes to protocol buffers
- `pglogical_output`
 - Plugin that `pglogical` uses
 - “Native” protocol
 - JSON protocol



pglogical



pglogical





Logical Replication

- Target node is writeable
 - Allows temp tables
 - Allows different indexes
 - Allows different security
 - Allows (some) data transformation
- Selective Replication
 - Can replicate subset of database
- Cross-version



pglogical

- Replicates transactions in commit order
- Selective Replication
- Online Upgrade
- Data Transport
 - Data integration
 - Streaming changes to analytical database
 - Master configuration data management
 - ...
- Optionally semi-synchronous



Replication Sets

- Replication is defined in terms of groups (sets) of tables, rather than individual tables
 - Need to be defined on each upstream node
- Table is not replicated until added to a set
- Tables may be defined in more than one set, but changes for the table will only be sent **once** to each subscriber
- Predefined sets, “default”, “default_insert_only”, “ddl_sql”



Replication Actions

- By default replication sets replicate all actions
 - INSERT, UPDATE, DELETE, TRUNCATE
- It's possible to filter actions for given replication set
- Useful for data aggregation, data warehousing etc.




Installation

- Extension
 - All configuration in database
- Provider + Subscriber
 - `CREATE EXTENSION pglogical;`
 - `create_node(name, dsn)`
- Provider
 - `replication_set_add_table(set_name, table_name)`
- Subscriber
 - `create_subscription(name, provider_dsn, ...)`




Example setup





Performance





DDL Replication

- Initial schema either fully synchronized or not at all
- The DDL commands are not automatically replicated yet
- `pglogical.replicate_ddl_command(command [, replication_sets])`



Other caveats

- Sequences not replicated yet
 - Fixed in upcoming 1.1
- Big transactions
- Does not play well with physical replication yet
 - Failover
 - Fixed in PostgreSQL 9.6



Future plans

- Make it available in PostgreSQL out of the box
 - 9.7?
- Data transformation hooks
- Filtering by expression
- Push replication
- Integrate BDR features
 - Multi-master
 - Transparent DDL replication



Questions?

- petr@2ndquadrant.com
- <http://2ndquadrant.com/pglogical/>