

# Introduction to PostgreSQL

Don't Panic!

Federico Campoli

Brighton PostgreSQL Users Group

13 June 2016



# Table of contents

- 1 An elephant never forgets
- 2 A quick look to a powerful tool
- 3 NOSQL on Acid
- 4 Wrap up

# Table of contents

- 1 An elephant never forgets
- 2 A quick look to a powerful tool
- 3 NOSQL on Acid
- 4 Wrap up

# An elephant never forgets



PostgreSQL

# PostgreSQL, an history of excellence

- Created at Berkeley in 1982 by database's legend Prof. Stonebraker
- In the 1994 Andrew Yu and Jolly Chen added the SQL interpreter
- In the 1996 becomes an Open Source project.
- The project's name changes in PostgreSQL

# PostgreSQL, an history of excellence

- Created at Berkeley in 1982 by database's legend Prof. Stonebraker
- In the 1994 Andrew Yu and Jolly Chen added the SQL interpreter
- In the 1996 becomes an Open Source project.
- The project's name changes in PostgreSQL
- Fully ACID compliant
- High performance in read/write with the MVCC
- Tablespaces

# PostgreSQL, an history of excellence

- Created at Berkeley in 1982 by database's legend Prof. Stonebraker
- In the 1994 Andrew Yu and Jolly Chen added the SQL interpreter
- In the 1996 becomes an Open Source project.
- The project's name changes in PostgreSQL
- Fully ACID compliant
- High performance in read/write with the MVCC
- Tablespaces
- Runs on almost any unix flavour
- From the version 8.0 is native on \*cough\* MS Windows \*cough\*
- HA with hot standby and streaming replication
- Heterogeneous federation

# PostgreSQL, an history of excellence

- Created at Berkeley in 1982 by database's legend Prof. Stonebraker
- In the 1994 Andrew Yu and Jolly Chen added the SQL interpreter
- In the 1996 becomes an Open Source project.
- The project's name changes in PostgreSQL
- Fully ACID compliant
- High performance in read/write with the MVCC
- Tablespaces
- Runs on almost any unix flavour
- From the version 8.0 is native on \*cough\* MS Windows \*cough\*
- HA with hot standby and streaming replication
- Heterogeneous federation
- Procedural languages (pl/pgsql, pl/python, pl/perl...)
- Support for NOSQL features like HSTORE and JSON



# Development

- Old ugly C language
- New development cycle starts usually in June
- New version released usually by the end of the year
- At least 4 LTS versions
- Can be extended using shared libraries
- Extensions (9.1+)
- BSD like license

# Limits

- Database size. No limits.
- Table size 32 TB
- Row size 1.6 TB
- Field size 1 GB
- Rows in table. No limits.
- Fields in table 250 - 1600 depending on data type.
- Tables in a database. No limits.

# Table of contents

- 1 An elephant never forgets
- 2 A quick look to a powerful tool
- 3 NOSQL on Acid
- 4 Wrap up

# A quick look to a powerful tool



Image by Hein Waschefort -

[http://commons.wikimedia.org/wiki/User:Hein\\_waschefort](http://commons.wikimedia.org/wiki/User:Hein_waschefort)

# Data types

PostgreSQL comes with an incredibly rich data type set.

# Data types

## Numerical

- smallint, integer, bigint
- decimal, numeric, user-specified precision, exact
- real, double precision, variable precision, inexact

# Data types

## Character

- character
- character varying with max size
- text, character varying

# Data types

## Binary

- bytea



# Data types

Alongside the general purpose data types PostgreSQL have some exotic types.

- Range (integers, date)
- Geometric (points, lines etc.)
- Network addresses
- XML
- JSON
- JSONB
- HSTORE (extension)

# UPSERT (9.5+)

```
INSERT INTO t_table (i_id, v_value)
VALUES (1, 'fake value'), (2, 'another fake value')
ON CONFLICT (i_id) DO UPDATE SET v_value = EXCLUDED.v_value;
```

# Row Level Security (9.5+)

Row Level Security, allows security "policies" filtering rows per database user.

# Big Data! (9.5+)

## BIG DATA!

- BRIN - Block Range Indices
- IMPORT FOREIGN SCHEMA, Federation with steroids
- TABLESAMPLE

# GROUPING SETS

GROUPING SETS (shameless copied from the on line manual)

The data selected by the FROM and WHERE clauses is grouped separately by each specified grouping set, aggregates computed for each group just as for simple GROUP BY clauses, and then the results returned. For example:

```
=> SELECT * FROM items_sold;
```

brand	size	sales
Foo	L	10
Foo	M	20
Bar	M	15
Bar	L	5

(4 rows)

```
=> SELECT brand, size, sum(sales) FROM items_sold GROUP BY GROUPING SETS ((
    brand), (size), ());
```

brand	size	sum
Foo		30
Bar		20
	L	15
	M	35
		50

(5 rows)

# The future

## PostgreSQL 9.6

### Currently in beta

- Parallel sequential scans, joins and aggregates
- Push down for the postgresql foreign data wrapper
- Full text search for phrases, YAY!
- Multiple synchronous standby servers
- Snapshot too old!

# Table of contents

- 1 An elephant never forgets
- 2 A quick look to a powerful tool
- 3 **NOSQL on Acid**
- 4 Wrap up

# NOSQL on Acid





# JSON

## JSON - JavaScript Object Notation

- The version 9.2 adds JSON as native data type
- The version 9.3 adds the support functions for JSON
- JSON is stored as text
- JSON is parsed and validated on the fly
- The 9.4 adds JSONB (binary) data type
- The 9.5 improves JSONB

# JSON

## JSON - Examples

### From record to JSON

```
postgres=# SELECT row_to_json(ROW(1,'foo'));
           row_to_json
-----
{"f1":1,"f2":"foo"}
(1 row)
```

### Expanding JSON into key to value elements

```
postgres=# SELECT * from json_each('{"a":"foo", "b":"bar"}');

 key | value
-----+-----
  a   | "foo"
  b   | "bar"
(2 rows)
```

# JSONB

Because JSON is parsed and validated on the fly it could be a bottleneck.

The new JSONB introduced with PostgreSQL 9.4 is parsed, validated and transformed at insert/update's time. The access is then faster than the plain JSON but the storage cost can be higher.

The functions available for JSON are also available in the JSONB flavour.

# Some numbers

Let's create three tables with text,json and jsonb type fields.  
Each record contains the same json element generated on  
<http://beta.json-generator.com/4kwCt-fwg>

```
[ {
  "_id": "56891aba27402de7f551bc91",
  "index": 0,
  "guid": "b9345045-1222-4f71-9540-6ed7c8d2ccae",
  "isActive": false,
  .....
  3,
  {
    "id": 1,
    "name": "Bridgett Shaw"
  }
],
  "greeting": "Hello, Johnston! You have 8 unread messages.",
  "favoriteFruit": "apple"
}
```

# Some numbers

```
DROP TABLE IF EXISTS t_json ;  
DROP TABLE IF EXISTS t_jsonb ;  
DROP TABLE IF EXISTS t_text ;
```

```
CREATE TABLE t_json as  
SELECT  
'<JSON ELEMENT>'::json as js_value  
FROM  
generate_series(1,100000);  
Query returned successfully: 100000 rows affected, 14504 ms execution time.
```

```
CREATE TABLE t_text as  
SELECT  
'<JSON ELEMENT>'::text as t_value  
FROM  
generate_series(1,100000);  
Query returned successfully: 100000 rows affected, 14330 ms execution time.
```

```
CREATE TABLE t_jsonb as  
SELECT  
'<JSON ELEMENT>'::jsonb as jsb_value  
FROM  
generate_series(1,100000);  
Query returned successfully: 100000 rows affected, 14060 ms execution time.
```

# Table size

```
SELECT
    pg_size_pretty(pg_total_relation_size(oid)),
    relname
FROM
    pg_class
WHERE
    relname LIKE 't\_%',
;
```

```
pg_size_pretty | relname
-----+-----
270 MB         | t_json
322 MB         | t_jsonb
270 MB         | t_text
(3 rows)
```

# Sequential scans

## TEXT

```
EXPLAIN (BUFFERS, ANALYZE) SELECT * FROM t_text;
```

```
Seq Scan on t_text (cost=0.00..1637.00 rows=100000 width=18) (actual time  
=0.016..17.624 rows=100000 loops=1)  
  Buffers: shared hit=637  
  Planning time: 0.040 ms  
  Execution time: 28.967 ms  
(4 rows)
```

# Sequential scans

## JSON

```
EXPLAIN (BUFFERS, ANALYZE) SELECT * FROM t_json;
```

```
Seq Scan on t_json (cost=0.00..1637.09 rows=100009 width=32) (actual time  
=0.018..15.443 rows=100000 loops=1)  
  Buffers: shared hit=637  
  Planning time: 0.045 ms  
  Execution time: 25.268 ms  
(4 rows)
```



# Sequential scans

## JSONB

```
EXPLAIN (BUFFERS, ANALYZE) SELECT * FROM t_jsonb;
```

```
Seq Scan on t_jsonb (cost=0.00..1637.00 rows=100000 width=18) (actual time  
=0.015..18.943 rows=100000 loops=1)  
  Buffers: shared hit=637  
  Planning time: 0.043 ms  
  Execution time: 31.072 ms  
(4 rows)
```

# Sequential scan with json access

## TEXT

```
EXPLAIN (BUFFERS, ANALYZE) SELECT t_value::json->'index' FROM t_text;
```

```
Seq Scan on t_text (cost=0.00..2387.00 rows=100000 width=18) (actual time  
=0.159..7748.381 rows=100000 loops=1)  
  Buffers: shared hit=401729  
    Planning time: 0.028 ms  
    Execution time: 7760.263 ms  
(4 rows)
```

# Sequential scan with json access

## JSON

```
EXPLAIN (BUFFERS, ANALYZE) SELECT js_value->'index' FROM t_json;
```

```
Seq Scan on t_json (cost=0.00..1887.11 rows=100009 width=32) (actual time  
=0.254..5787.267 rows=100000 loops=1)  
  Buffers: shared hit=401730  
  Planning time: 0.044 ms  
  Execution time: 5798.153 ms  
(4 rows)
```

# Sequential scan with json access

## JSONB

```
EXPLAIN (BUFFERS, ANALYZE) SELECT jsb_value->'index' FROM t_jsonb;
```

```
Seq Scan on t_jsonb (cost=0.00..1887.00 rows=100000 width=18) (actual time  
=0.138..1678.222 rows=100000 loops=1)  
  Buffers: shared hit=421729  
  Planning time: 0.048 ms  
  Execution time: 1688.752 ms  
(4 rows)
```

# Table of contents

- 1 An elephant never forgets
- 2 A quick look to a powerful tool
- 3 NOSQL on Acid
- 4 Wrap up**

# Wrap up

PostgreSQL is a powerful RDBMS with dozens of features and capabilities.

Choosing the correctly what to use can be tricky.

The lack of horizontal scalability in PostgreSQL can be a problem for large data sets. However, an interesting project for a distributed cluster is PostgreSQL XL - <http://www.postgres-xl.org/>

Another cool project is CitusDB, a powerful extension to add to PostgreSQL horizontal scale capabilities.

CitusDB recently un-forked from PostgreSQL becoming an open source extension. Yay!

# Wrap up

- Schema less data are useful. They are flexible and powerful.
- Never forget PostgreSQL is a RDBMS
- Get a DBA on board

# Questions

Questions?



# Contacts

- Twitter: 4thdoctor\_scarf
- Personal blog: <http://www.pgdba.co.uk>
- PostgreSQL Book:  
<http://www.slideshare.net/FedericoCampoli/postgresql-dba-01>
- Brighton PostgreSQL Meetup:  
<http://www.meetup.com/Brighton-PostgreSQL-Meetup/>

# License and copyright

This presentation is licensed under the terms of the Creative Commons Attribution NonCommercial ShareAlike 4.0



<http://creativecommons.org/licenses/by-nc-sa/4.0/>

- The cheetah photo is copyright by Hein Waschefort - [http://commons.wikimedia.org/wiki/User:Hein\\_waschefort](http://commons.wikimedia.org/wiki/User:Hein_waschefort)
- The elephant logos are copyright of the PostgreSQL Global Development Group - <http://www.postgresql.org/>

# Introduction to PostgreSQL

Don't Panic!

Federico Campoli

Brighton PostgreSQL Users Group

13 June 2016

