

PostgreSQL Streaming Replication Cheat Sheet

Master setup	Standby setup
<p>Primary GUC's (also needs postgresql restart)</p> <p>wal_level = hot_standby – required for hot-standby.</p> <p>max_wal_senders = 6 – max number of concurrent replication connections (must be > 0).</p> <p>Secondary GUC's</p> <p>wal_keep_segments = 256 – (4GB on disk) or set more for high-write and huge databases.</p> <p>wal_sender_timeout = 60s – ok by default.</p>	<p>Primary GUC's</p> <p>hot_standby = on – if enabled standby can accept read-only queries.</p> <p>Secondary GUC's</p> <p>max_standby_streaming_delay = 300s – set more if recovery conflicts occurs.</p> <p>wal_receiver_status_interval = 10s – ok by default.</p> <p>hot_standby_feedback = on – enable it anyway.</p> <p>wal_receiver_timeout = 60s – ok by default.</p>
Create dedicated role for replication (optional, need reload)	Standby recovery.conf (streaming replication)
<p>pg_hba.conf format</p> <p>host replication repluser standby_ip/netmask authtype (trust, md5, etc.)</p> <p>Add role with psql and reload configuration:</p> <pre>CREATE ROLE repluser WITH LOGIN REPLICATION PASSWORD 'lovesexgod'; SELECT pg_reload_conf();</pre>	<pre>primary_conninfo = 'host=master port=... user=... password=...' standby_mode = 'on' – start as standby and don't stop recovery. trigger_file = '/path/to/file' – create this file to switch postgres from standby mode to normal. recovery_min_apply_delay = 5min – delay recovery by a fixed period of time (since 9.4).</pre>
pg_basebackup	Standby recovery.conf (point-in-time-recovery)
<p>Main options</p> <p>-h master, -p port, -U user, -D destdir – (destdir must have 700 perms).</p> <p>Auxiliary options</p> <p>-c fast spread – use 'fast' for start basebackup as soon as possible, 'spread' for minimize load.</p> <p>-X stream fetch – add WAL archives into backup. The 'stream' uses separate streaming connection (since 9.2).</p> <p>-R – create minimal recovery.conf (since 9.3).</p> <p>-r – limit network bandwidth in kB/s, or use suffix "k" or "M" (since 9.4).</p> <p>--xlogdir=dir – set new pg_xlog location (since 9.4).</p> <p>-T olddir=newdir – set new locations for tablespaces (since 9.4).</p> <p>-P – show progress.</p>	<pre>restore_command = 'cp /path/to/archives/%f "%p"' – command for restore WAL. archive_cleanup_command = 'pg_archivecleanup /archivedir %r' – clean up old WAL archives. recovery_target = 'immediate' – stop recover when consistent state is reached. recovery_target_name = 'deploy' – recover to point created by pg_create_restore_point(). recovery_target_time = '2016-01-06 18:11:54.840563' – recover to timestamp. recovery_target_xid = 1234567 – recover to transaction ID. recovery_target_inclusive = true – recover specified target (true) or stop before target (false). recovery_target_timeline = 'latest' – recover to specified timeline number. recovery_end_command = 'cmd' – run cmd after recovery. recovery_target_action = pause promote shutdown – take action after recovery (since 9.5).</pre>
Synchronous replication	Replication slots
<p>Master postgresql.conf</p> <p>synchronous_standby_names = 'main'</p> <p>Standby recovery.conf</p> <p>primary_conninfo = '... application_name = main ...'</p>	<p>Master postgresql.conf</p> <p>max_worker_processes = max_replication_slots = 6 – max number of replication slots.</p> <p>Create slot on master with psql</p> <pre>SELECT pg_create_physical_replication_slot('main');</pre> <p>Standby recovery.conf</p> <p>primary_slot_name = 'main'</p>
Monitoring on master	Monitoring on standby
<p>SELECT client_addr, pg_xlog_location_diff(pg_current_xlog_location(),sent_location) as pending from pg_stat_replication; – get pending (not sent) amount of local WAL on master. Big values (more than 1GB) is bad and mean that the master under heavy load.</p> <p>SELECT client_addr, pg_xlog_location_diff(sent_location,replay_location) as lag from pg_stat_replication; – get sent but not replayed amount of WAL on standby. Big values (more than 1GB) is bad and mean that the standby under heavy load or has network problems.</p>	<p>SELECT now() - pg_last_xact_replay_timestamp(); – get lag in seconds;</p>